SPSC Bounded Queue

Ditch the lock, speed the Queue!

Sarthak Sehgal

C++ Software Engineer

- Working at a high frequency options market making firm
- Interested in finance, low level programming, and C++ under the hood
- sartech.substack.com
- LinkedIn://sarthaksehgal99

Setup

- Exactly one producer and one consumer
- Fixed capacity
- Producer and consumer threads are pinned to separate physical cores and continuously poll for data

Setup

- Exactly one producer and one consumer
- Fixed capacity
- Producer and consumer threads are pinned to separate physical cores and continuously poll for data

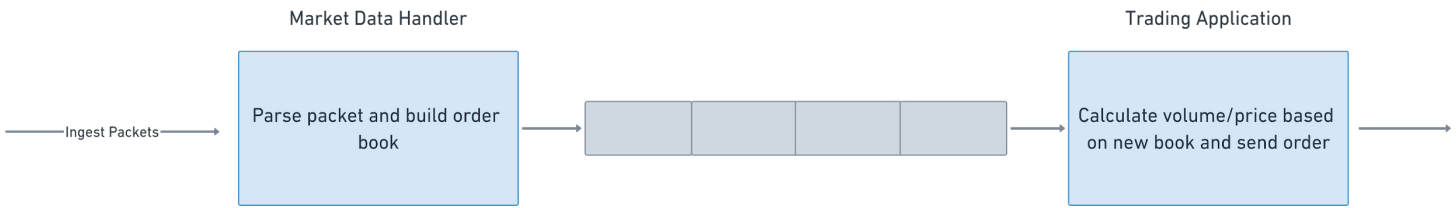

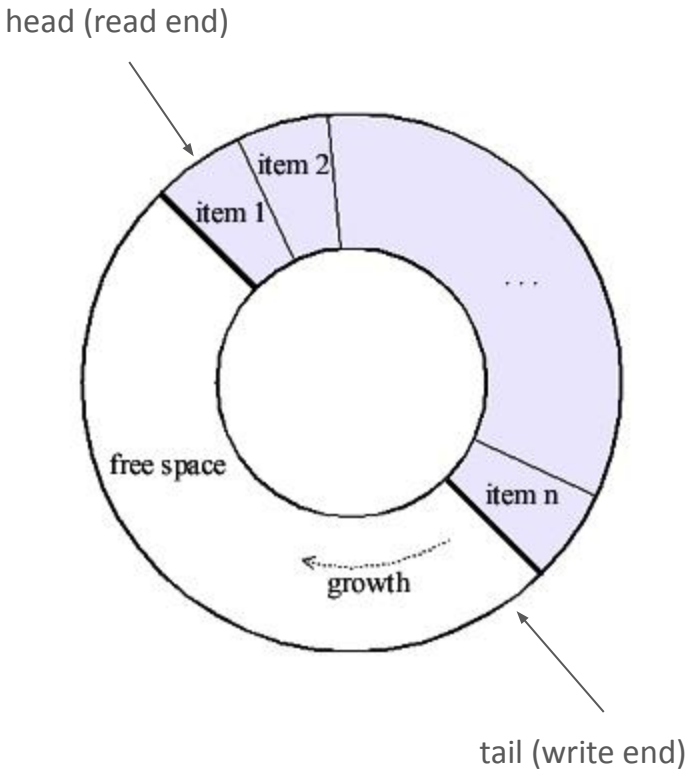
head (read end)
item 1
item 2
. . .
item n
free space
growth
tail (write end)

v1: Good old mutex

```
bool push(T const& item)
{
  std::lock_guard<std::mutex> lock(mtx);
  // access to buffer is safe
  ...
}

bool consume_one(auto&& func)
{
  std::lock_guard<std::mutex> lock(mtx);
  ...
}
```

v1: Good old mutex

```cpp
bool push(T const& item)
{
  std::lock_guard<std::mutex> lock(mtx);
  // access to buffer is safe
  ...
}

bool consume_one(auto&& func)
{
  std::lock_guard<std::mutex> lock(mtx);
  ...
}
```

blocking, expensive system calls

v2: Using atomics

```
bool push(const T& item)
{
  auto currTail = tail.load();
  auto nextTail = currTail+1 == capacity ? 0 : currTail+1;
  if (nextTail == head.load())
      return false;
  new (buffer + currTail) T(item);
  tail.store(nextTail);
  return true;
}
```

```
bool consume_one(auto&& func)
{
  auto currHead = head.load();
  if (currHead == tail.load())
      return false;
  T* elem = reinterpret_cast<T*>(buffer+currHead);
  func(*elem);
  elem->~T();
  auto nextHead = currHead+1 == capacity ? 0 : currHead+1;
  head.store(nextHead);
  return true;
}
```

atomics are used to *synchronize access* to the shared memory

v2: Using atomics

```
bool push(const T& item)
{
  auto currTail = tail.load();
  auto nextTail = currTail+1 == capacity ? 0 : currTail+1;
  if (nextTail == head.load())
      return false;
  new (buffer + currTail) T(item);
  tail.store(nextTail);
  return true;
}
```

```
bool consume_one(auto&& func)
{
  auto currHead = head.load();
  if (currHead == tail.load())
      return false;
  T* elem = reinterpret_cast<T*>(buffer+currHead);
  func(*elem);
  elem->~T();
  auto nextHead = currHead+1 == capacity ? 0 : currHead+1;
  head.store(nextHead);
  return true;
}
```

atomics are used to *synchronize access* to the shared memory

non-blocking ✅ lock-free ✅ strict memory ordering ⚠️

Out of order execution

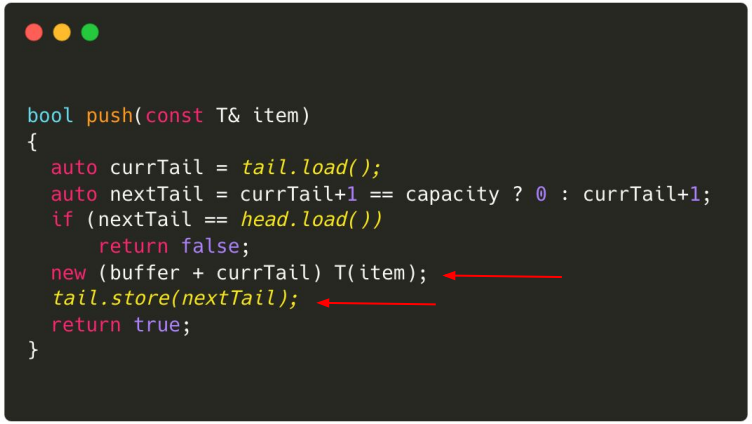

```
bool push(const T& item)
{
  auto currTail = tail.load();
  auto nextTail = currTail+1 == capacity ? 0 : currTail+1;
  if (nextTail == head.load())
      return false;
  new (buffer + currTail) T(item);
  tail.store(nextTail);
  return true;
}
```

Out of order execution

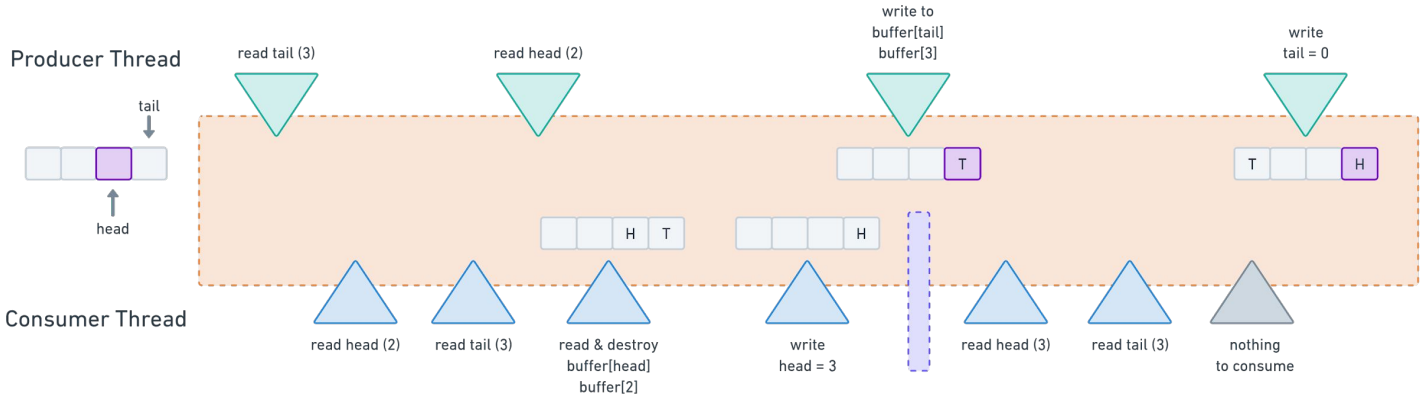

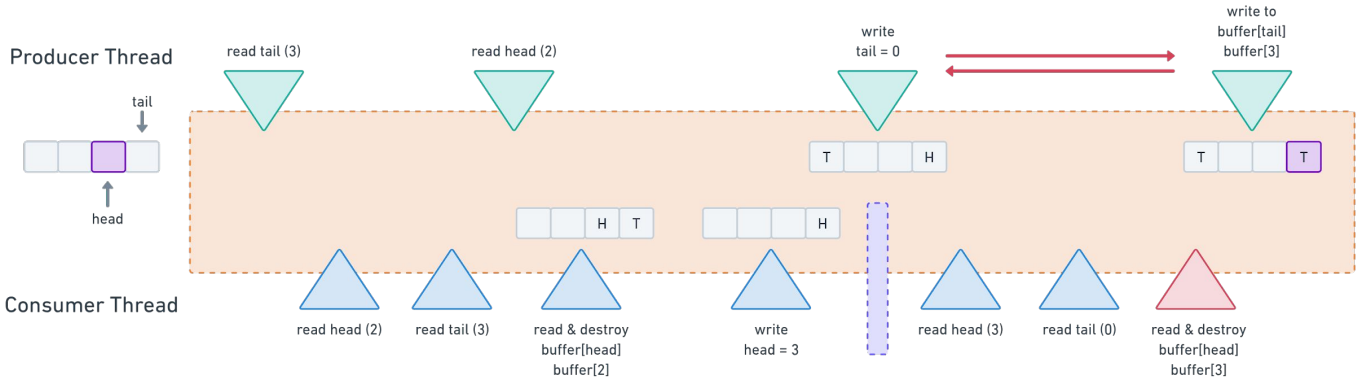

Memory Ordering

relaxed: only atomicity, no ordering constraints

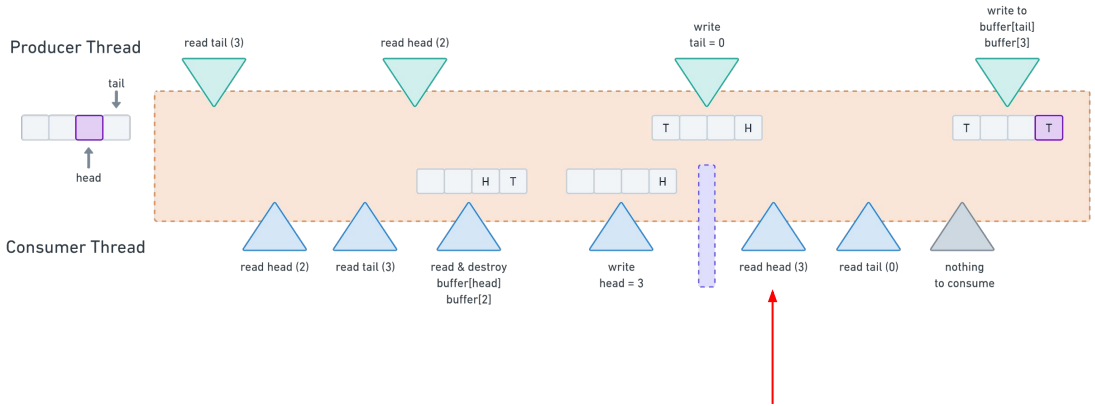

v2.1: Optimized memory ordering

```cpp
bool push(const T& item)
{
  auto currTail = tail.load(std::memory_order_relaxed);
  auto nextTail = currTail+1 == capacity ? 0 : currTail+1;
  if (nextTail == head.load(std::memory_order_acquire))
      return false;
  new (buffer + currTail) T(item);
  tail.store(nextTail, std::memory_order_release);
  return true;
}
```

```cpp
bool consume_one(auto&& func)
{
  auto currHead = head.load(std::memory_order_relaxed);
  if (currHead == tail.load(std::memory_order_acquire))
      return false;
  T* elem = reinterpret_cast<T*>(buffer+currHead);
  func(*elem);
  elem->~T();
  auto nextHead = currHead+1 == capacity ? 0 : currHead+1;
  head.store(nextHead, std::memory_order_release);
  return true;
}
```

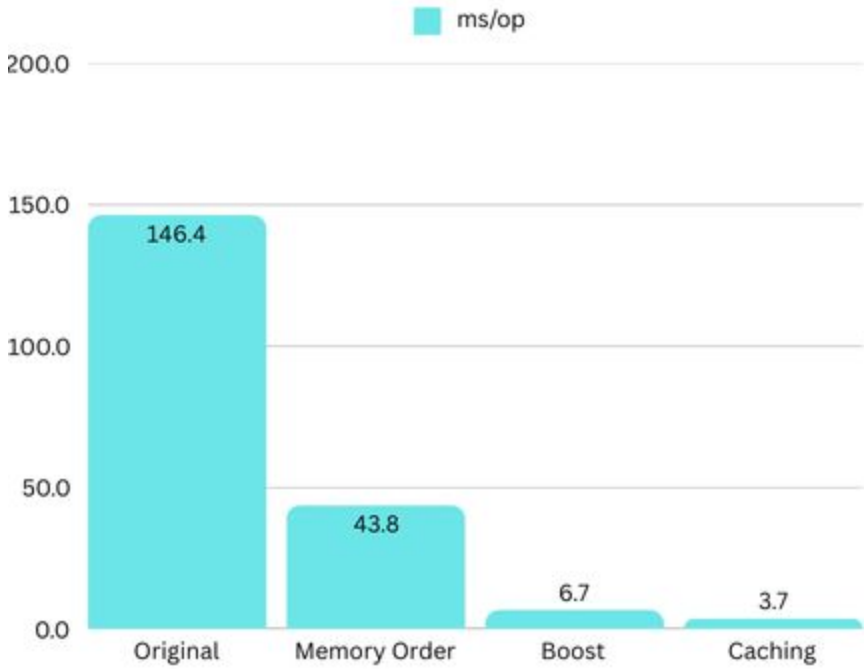
ms/op
200.0
150.0
100.0
50.0
0.0
146.4
43.8
6.7
3.7
Original
Memory Order
Boost
Caching

False Sharing

```
Element* buffer;
std::size_t const capacity;
std::atomic<std::size_t> head = 0;
std::atomic<std::size_t> tail = 0;
```

False Sharing

```
std::size_t const capacity;
Element* buffer;
alignas(std::hardware_destructive_interference_size) std::atomic<std::size_t> head = 0;
alignas(std::hardware_destructive_interference_size) std::atomic<std::size_t> tail = 0;
```

Minimum offset between two objects to avoid false sharing

Since C++17

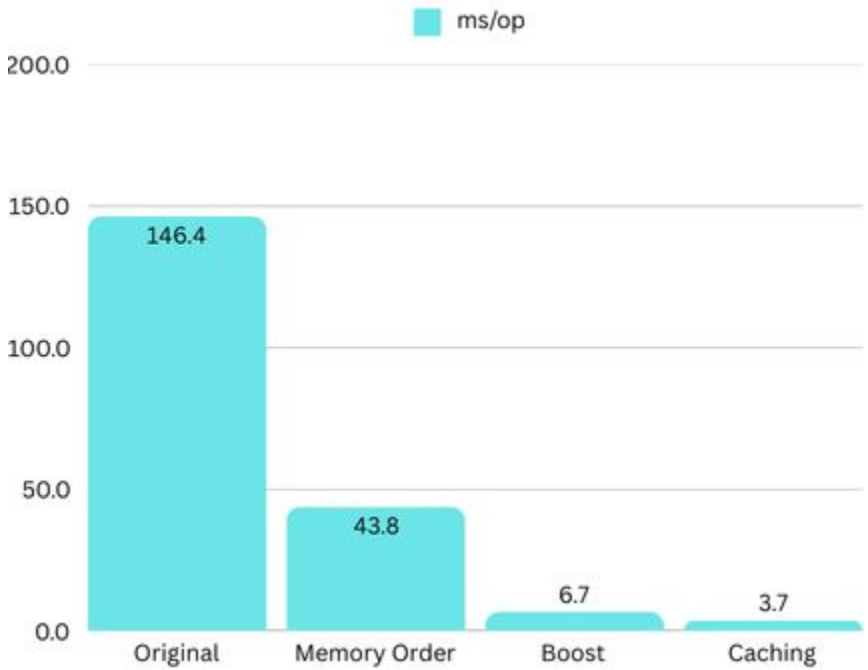

Interested to learn more? Hop over to my poster booth tomorrow

Memory Ordering

release: no reads/writes in current thread can be reordered after the *store*

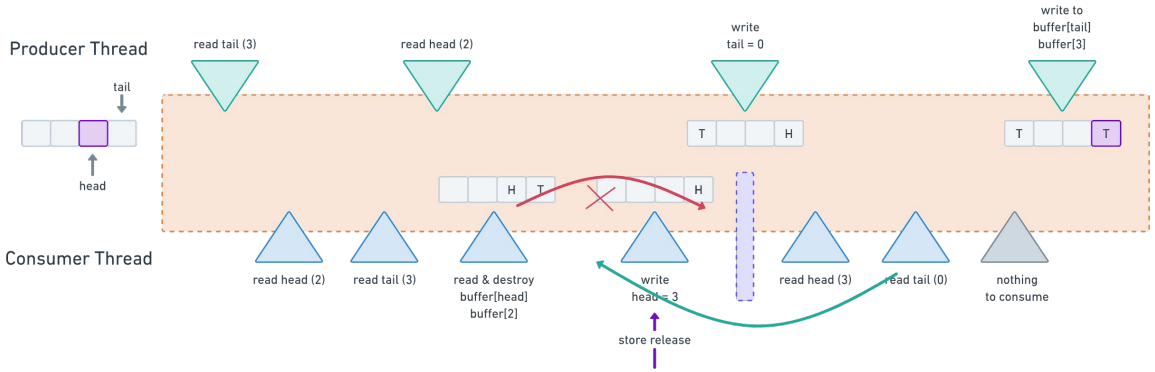

Memory Ordering

acquire: no reads/writes in current thread can be reordered before the *load*

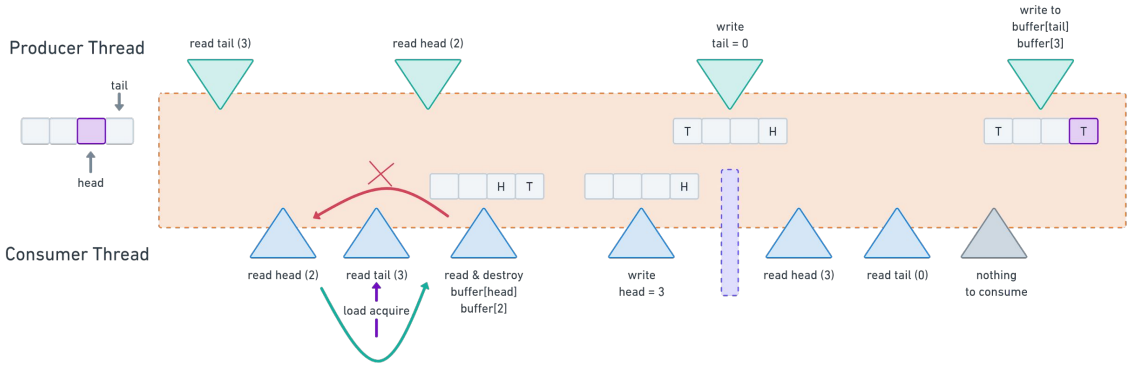

All writes in other threads that release the same atomic variable are visible in the current thread