

C++ONLINE

ROI BARKAN

TALK:

MORE RANGES PLEASE

[Slides](#)



2025

Hi, I'm Roi

- Roi Barkan (he/him) - רועי ברקן
- I live in Tel Aviv 🏠
- C++ developer since 2000
- SVP Development & Technologies @ Istra Research
 - Finance, Low Latency, in Israel
 - careers@istraresearch.com
- Always happy to learn and explore
 - Please - ask questions, make comments



[Slides](#)

Outline

- Ranges
 - Introduction
 - Strengths, core ideas
 - Review the views
- Libraries
 - What
 - Why
 - How
- Rabbits
- Summary

C++ Ranges

Ranges is a Breakthrough Library

- One of C++20 big-four features
- Rests on decades of existing libraries and experience
 - C++98 iterator-based algorithms
 - Fundamentals of functional / vectoric languages (APL, BQN, R, Julia, NumPy) [Conor Hoekstra](#)
 - Libraries of similar languages (D, Rust, Java) [Barry Revzin](#).
- Main Innovation - Composability
 - Many algorithms take ranges as input and return ranges as output
 - Opposed to in-place or output-iterator nature of C++98 algorithms
 - Range Adaptors - algorithms encapsulated as 'lazy ranges' (views)
 - Algorithms as composable objects - 'expression templates'
 - Projections - unary transformations of the ranges we inspect.

Composability of Ranges

- Chaining algorithms due to range arguments and results
`ranges::reverse(ranges::search(str, "abc"sv)); godbolt`
- Views as composable lazy ranges
`str | views::split(' ') | views::take(2); godbolt`
- Views have a value/algorithm duality
`auto square_evens =
 views::filter([](auto x) { return int(x) % 2 == 0; }) |
 views::transform([](auto x) { return x * x; }); godbolt`

The Views in the Standard (C++20/23*/26**)

- Factories/Generators: `empty`, `single`, `iota`, `repeat*`, `(std::generator*)`
- Rank preserving: `all`, `filter`, `transform`, `take{_while}`,
`drop{_while}`, `reverse`, `stride*`, `adjacent_transform*`,
`cache_latest**`, `(counted)`
- Rank preserving - variadic→tuples: `zip*`, `cartesian_product*`
- Rank decreasing - tuples: `elements`, `keys`, `values`
- Rank decreasing - variadic: `zip_transform*`, `concat**`
- Rank decreasing - ranges: `join{_with*}`
- Rank increasing - tuples: `enumerate*`, `adjacent*`
- Rank increasing - ranges: `{lazy_}split`, `slide*`, `chunk{_by}*`
- Committee plan for C++26 is in [P2760](#)

Factories / Generators

```
namespace stdv = std::views;
```

```
stdv::empty<char>
```

```
//=> []
```

```
stdv::single('+')
```

```
//=> ['+']
```

```
stdv::iota(2,5)
```

```
//=> [2, 3, 4]
```

```
stdv::repeat(0.3,3)
```

```
//=> [0.3, 0.3, 0.3]
```

[godbolt](https://godbolt.com)

Rank Preserving - 1/2

```
auto not5 = [](int i){return i != 5;};  
auto mult2 = [](int i){return i * 2;};  
auto iota2_10 = stdv::iota(2,10);
```

```
iota2_10 | stdv::all           //=> [2, 3, 4, 5, 6, 7, 8, 9]  
iota2_10 | stdv::filter(not5)  //=> [2, 3, 4, 6, 7, 8, 9]  
iota2_10 | stdv::transform(mult2) //=> [4, 6, 8, 10, 12, 14, 16, 18]  
iota2_10 | stdv::take(6)       //=> [2, 3, 4, 5, 6, 7]  
iota2_10 | stdv::drop(6)       //=> [8, 9]
```

[godbolt](https://godbolt.com)

Rank Preserving - 2/2

```
auto not5 = [](int i){return i != 5;};
```

```
auto iota2_10 = stdv::iota(2,10);
```

```
iota2_10 | stdv::take_while(not5)           //=> [2, 3, 4]
```

```
iota2_10 | stdv::drop_while(not5)          //=> [5, 6, 7, 8, 9]
```

```
iota2_10 | stdv::reverse                   //=> [9, 8, 7, 6, 5, 4, 3, 2]
```

```
iota2_10 | stdv::stride(3)                 //=> [2, 5, 8]
```

```
iota2_10 | stdv::adjacent_transform<2>(std::plus{})  
                                                //=> [5, 7, 9, 11, 13, 15, 17]
```

```
iota2_10 | stdv::cache_latest              //=> [2, 3, 4, 5, 6, 7, 8, 9]
```

[godbolt](https://godbolt.com)

Rank Preserving - Variadic \Rightarrow Tuples

```
auto iota2_7 = stdv::iota(2,7);
```

```
auto iota2_4 = stdv::iota(2,4);
```

```
auto iota6_9 = stdv::iota(6,9);
```

```
stdv::zip(iota2_7, iota6_9) //=> [(2, 6), (3, 7), (4, 8)]
```

```
stdv::cartesian_product(iota2_4, iota6_9) //=> [(2, 6), (2, 7), (2, 8),  
                                                (3, 6), (3, 7), (3, 8)]
```

[godbolt](https://godbolt.com)

Rank Decreasing - Tuples

```
auto the_zip = std::zip(iota2_7, iota6_9, "abcdef"sv);  
                //=> [(2, 6, 'a'), (3, 7, 'b'), (4, 8, 'c')]
```

```
the_zip | std::keys           //=> [2, 3, 4]
```

```
the_zip | std::values         //=> [6, 7, 8]
```

```
the_zip | std::elements<2>    //=> ['a', 'b', 'c']
```

Rank Decreasing - Variadic

```
auto iota2_7 = stdv::iota(2,7); auto iota6_9 = stdv::iota(6,9);  
  
stdv::zip_transform(iota2_7, iota6_9, std::multiplies{})  
                                     //=> [12, 21, 32]  
stdv::concat(iota2_7, iota6_9)      //=> [2, 3, 4, 5, 6, 6, 7, 8]
```

Rank Decreasing - Ranges

```
vector{"hey"sv, "C++"sv} | stdv::join
                               //=> ['h', 'e', 'y', 'C', '+', '+']

(vector{"hey"sv, "C++"sv} | stdv::join_with(':', ' '))
                               //=> ['h', 'e', 'y', ':', 'C', '+', '+']
```

Rank Increasing - Tuples

```
"hey"sv | stdv::enumerate
```

```
//=> [(0, 'h'), (1, 'e'), (2, 'y')]
```

```
"hello"sv | stdv::adjacent<3>
```

```
//=> [('h', 'e', 'l'), ('e', 'l', 'l'), ('l', 'l', 'o')]
```

Rank Increasing - Ranges

```
"hey C++"sv | stdv::split(' ') //=> [['h', 'e', 'y'], ['C', '+', '+']]
```

```
"hey C++"sv | stdv::lazy_split(' ')  
//=> [['h', 'e', 'y'], ['C', '+', '+']]
```

```
"hello"sv | stdv::slide(3) //=> [['h', 'e', 'l'], ['e', 'l', 'l'],  
                                ['l', 'l', 'o']]
```

```
"hey C++"sv | stdv::chunk(3) //=> [['h', 'e', 'y'], [' ', 'C', '+'],  
                                ['+']]
```

```
"hello C++"sv | stdv::chunk_by(equal_to{})  
//=> [['h'], ['e'], ['l', 'l'], ['o'], [' ', 'C'], ['+', '+']]
```

[godbolt](https://godbolt.com)

Recap: Views in the STL (C++20/23*/26**)

- Factories/Generators: `empty`, `single`, `iota`, `repeat*`, `(std::generator*)`
- Rank preserving: `all`, `filter`, `transform`, `take{ _while }`,
`drop{ _while }`, `reverse`, `stride*`, `adjacent_transform*`,
`cache_latest**`, `(counted)`
- Rank preserving - variadic→tuples: `zip*`, `cartesian_product*`
- Rank decreasing - tuples: `elements`, `keys`, `values`
- Rank decreasing - variadic: `zip_transform*`, `concat**`
- Rank decreasing - ranges: `join{ _with* }`
- Rank increasing - tuples: `enumerate*`, `adjacent*`
- Rank increasing - ranges: `{lazy_}split`, `slide*`, `chunk{ _by }*`
- Committee plan for C++26 is in [P2760](#)

Libraries

What is a Library


- Code for coders
- Self contained reusable software
- Abstraction layer
- Language within a programming language
 - Domain specific language (DSL)
- Building blocks
- Simple interface (API) with tricky implementation
- Vocabulary
- ...

Clip

SICP

link

**Abelson,
Sussman,
Sussman
1984**



Why Should We Write Libraries

- Reuse: don't repeat yourself (DRY)
- Vocabulary: raise the level of discussion
- Building blocks:
 - The higher we go - change is more likely (close to users)
 - The lower we go - more stable, potentially robust.
- Abstraction layers:
 - Complexity is safer when it's well encapsulated
 - Lower level likely better defined (pre/post-conditions)
 - Lower level is typically more testable.
- Libraries can be stacked / composed.



Sean Parent's Vision

A (New) Possible Future

- A large library of proven generic components
- A small number of non-Turing complete declarative *forms* for assembling the generic components

How to Write Libraries

- Start with algorithms ([Christopher Di Bella](#), [Eric Niebler](#))
 - Write the code that you need
 - Generalize, distill, *simplify*
 - Concepts will emerge from the *generic* attributes you find.
- Pay attention to the API - the abstraction layer
 - Easy to use, hard to misuse ([Ben Deane](#))
 - Consider preconditions and postconditions
 - Return every potentially useful piece of information (Stepanov [Law of Useful Return](#)).
- Aim for *regular, algebraic* data types
 - Value semantic objects typically behave well in the language
 - Do you have a *group*?, a *monoid*?, a *monad*?

Finding Algorithms

- Algorithms are everywhere
 - Not just in CS textbooks
- Can be found in many places
 - Other languages
 - Other libraries
 - Research papers
 - Your own codebase

Clip

**Stepanov
on
Science**

link

**STL and its Design
Principles
2002**



More From Stepanov



[link](#)

More Ranges

First Inspiration - Hondt Method

- Credit to Tina Ulbrich - [How to Rangify Your Code](#)
- Hondt method - assigning parliament seats in a party-voting democracy

	party 1 votes: 110		party 2 votes: 85		party 3 votes: 35	
1	(1)	$110 / 1 = 110$	(2)	$85 / 1 = 85$	(6)	$35 / 1 = 35$
2	(3)	$110 / 2 = 55$	(4)	$85 / 2 = 42.5$		$35 / 2 = 17.5$
3	(5)	$110 / 3 = 36.66$	(7)	$85 / 3 = 28.33$		$35 / 3 = 11.66$
4		$110 / 4 = 27.5$		$85 / 4 = 21.25$		$35 / 4 = 8.75$
5		$110 / 5 = 22$		$85 / 5 = 17$		$35 / 5 = 7$
6		$110 / 6 = 18.33$		$85 / 6 = 14.16$		$35 / 6 = 5.83$
7		$110 / 7 = 15.71$		$85 / 7 = 12.14$		$35 / 7 = 5$
	seats: 3		seats: 3		seats: 1	

Basic Hondt Approach

	party 1 votes: 110		party 2 votes: 85		party 3 votes: 35	
1	(1)	$110 / 1 = 110$	(2)	$85 / 1 = 85$	(6)	$35 / 1 = 35$
2	(3)	$110 / 2 = 55$	(4)	$85 / 2 = 42.5$		$35 / 2 = 17.5$
3	(5)	$110 / 3 = 36.66$	(7)	$85 / 3 = 28.33$		$35 / 3 = 11.66$
4		$110 / 4 = 27.5$		$85 / 4 = 21.25$		$35 / 4 = 8.75$
5		$110 / 5 = 22$		$85 / 5 = 17$		$35 / 5 = 7$
6		$110 / 6 = 18.33$		$85 / 6 = 14.16$		$35 / 6 = 5.83$
7		$110 / 7 = 15.71$		$85 / 7 = 12.14$		$35 / 7 = 5$
		seats: 3			seats: 3	seats: 1

- Calculate cells and sort them

```
std::ranges::sort(proportional_votes, [] (const auto& rhs, const auto& lhs)
{
    return rhs.proportion > lhs.proportion;
});
proportional_votes.resize(total_number_of_seats);
```

- Key observations for P parties and S seats:
 - `sort()` requires eager evaluation and storage of S*P cells
 - `resize()` implies we are sorting too much - `nth_element()` requires O(SP) average steps
 - Each column is already sorted - we can `merge()`
 - Complexity can go down to O(SlogP)
 - `merge()` allow lazy evaluation - a range-view won't need pre-computation/allocation

Rabbit 1 - Views for Sorted Ranges

- Suggestion - views for `merge`, `set_union`, `set_intersection`, `set_{symmetric}_difference`
 - Most algorithms can benefit from multi-input implementations
 - Heap (`priority_queue`) is needed for efficient `set_union`, `merge`,
- STL contains several algorithms for sorted ranges: `{inplace}_merge`, `includes`, `set_{union,intersection,{symmetric}_difference}`
 - Also search algorithms: `{upper,lower}_bound`, `equal_range`, `(unique)`.
- All the operations are lazy in nature
- Ranges-v3 has views for `set_{union,intersection,{symmetric}_difference}` with 2 input ranges
- D-lang has `merge` and `multiWayMerge`.

Thoughts on Sorted Ranges (Rabbit 2)

- `equal_range`, `unique` can also be views - they have the right signature
- Binary search algorithm can become filtering views for random-access ranges
 - `take_until(value)`, `drop_until(value)` with $O(\log N)$ cost to cache the cut-off
 - Technically `take_until` mostly needs this to stay random-access.
 - `take_between(min, max)` can be more efficient than separate searches.
- If we'll have many algorithms specifically for sorted ranges, we might want a concept, which can also provide the comparator (`value_comp`)
 - D has [it](#), as well as [isSorted](#) and [assumeSorted](#) factories.
- Simple combinations can be good for vocabulary:

```
auto histogram =
    views::chunk_by(std::equals{ }) |
    views::transform([](const auto& rng) {
        return make_pair(begin(rng), distance(rng));
    });
```

Thoughts on Sorted Ranges (Rabbit 2)

- `equal_range`, `unique` can also be views - they have the right signature
- Search can also become filtering views for random-access ranges
 - `take_until(value)`, `drop_until(value)` with $O(\log N)$ cost to cache the cut-off
 - Technically `take_until` mostly needs this to stay random-access.
 - `take_between(min,max)` can be more efficient than separate searches.
- If we'll have many algorithms specifically for sorted ranges, we might want a concept, which can also provide the comparator (`value_comp`)
 - D has [it](#), as well as [isSorted](#) and [assumeSorted](#) factories.
- Simple combinations can be good for vocabulary:

```
auto histogram =  
    views::chunk_by(std::equals{}) |  
    views::transform([](const auto& rng) {  
        return make_pair(begin(rng), distance(rng));  
    });
```



Digression - Algorithm Selection

- Sometimes the same goal can be achieved in several ways
 - `ranges::distance` - returns the distance between the beginning and end of a range
 - `ranges::ssize` - returns a signed integer equal to the size of a range
 - `ssize` only works for ranges that have constant-time calculation (opt-out semantic concept); `distance` allows linear calculation. [Ben Deane recommends it.](#)
- The library uses concepts to constrain which ranges are applicable for which algorithm/view, and to know the best method of reaching the intended goal
- Before C++20 other mechanisms were used to achieve this goal - and with concepts we have a way to be more precise and more flexible where needed.

Breaking Sort Apart

- Inspired by the R standard library
- `sort()` is $O(N \log N)$ in two aspects:
 - Comparisons (+projections)
 - `iter_swap` operations.
- Sometimes `iter_swap` is much more expensive than comparison
 - For example - sorting rows of a table by one column
 - Projections stress the potential difference between the objects in question.
- Suggested algorithm: `order()` - generate the sorting *permutation*
 - Requires allocating N indexes, and $O(N \log N)$ comparisons
 - Then, *apply* the permutation with N `iter_swap` operations
 - Implementation is likely a one-liner using `iota()` and projections.

Rabbit 3 - Permutations

- Permutations allow more flexibility, not just intermediary step:
 - Can be reversed, to regain the original order
 - Can be lazily applied (be range views) - no swaps required.
- Iterator adapters that lazily apply permutations exist in [boost](#), [thrust](#).
- The standard library has other algorithms that deal with permutations.
- Potentially, there might be room for a `permutation`, or `permutation_of<Range>` concept.

Rabbit 3 - Permutations

- Permutations allow more flexibility, not just intermediary step:
 - Can be reversed, to regain the original order
 - Can be lazily applied (be range views) - no swaps required.
- Iterator adapters that lazily apply permutations exist in [boost](#), [thrust](#).
- The standard library has other algorithms that deal with permutations.
- Potentially, there might be room for a **permutation**, or **permutation_of<Range>** concept.



Summary

- The C++ ranges library is an exemplar of composability
 - Libraries can make software better, safer and cleaner
 - Potential libraries are all around us - it's not just the STL or written by others
 - Even C++ ranges might be improved/enhanced in novel ways
 - “Go catch rabbits”
-
- Thank you !!
 - Questions and comments are welcome
 - Thanks to Bryce Lebach for review and comments



[Slides](#)